

Highly accurate protein structure prediction with AlphaFold

AISC Healthcare Discussion Group

Willy Rempel

Saturday, July 24, 2021

Introduction



Willy Rempel

- HBS Sc Computer Science
- BSc Mathematics
- Research Associate, AISC
- seeking opportunities in the field

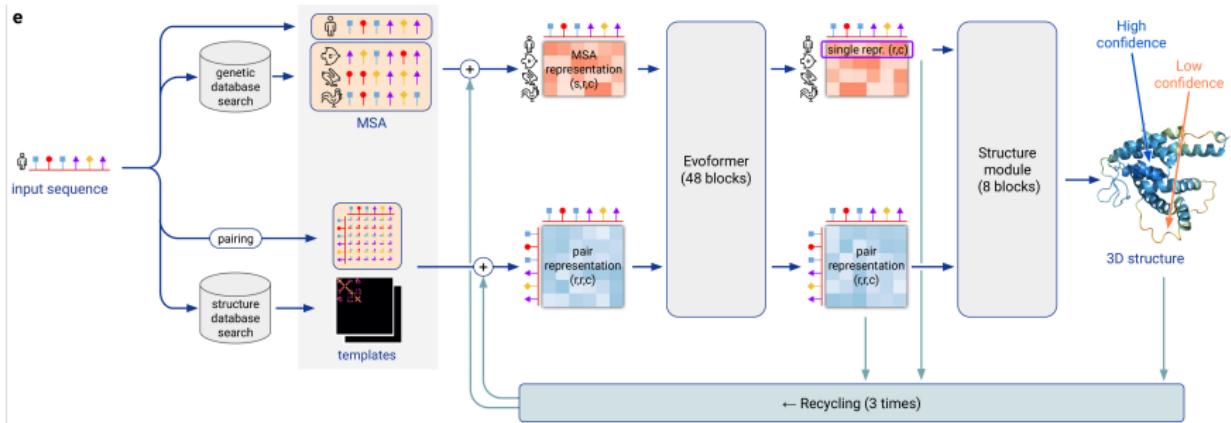
Introduction

Although all of the ideas in the model are doubtlessly clever, the main secret behind AlphaFold 2's success is the superb deep learning engineering. A close look at the model reveals an architecture with a large amount of small details that seem fundamental for the performance of the network. As we admire the end product, we should not turn a blind eye to the enormous budget, and the large team of full-time, handsomely paid engineers that made it possible. [3]

This, and many other tricks, are described in exhaustive detail in the Supplementary Information. A reduced subset has been analysed in a brief ablation study, but ultimately, how important are each of the minor details is anybody's guess. [3]

(above blog post is recommended reading)

Model Overview [1]



Initial Input: mmCIF or FASTA files

PDBx/mmCIF

```
loop_
_atom_site.group_PDB
_atom_site.id
_atom_site.type_symbol
_atom_site.label_atom_id
_atom_site.label_alt_id
_atom_site.label_comp_id
_atom_site.label_asym_id
_atom_site.label_entity_id
_atom_site.label_seq_id
_atom_site.pdbx_PDB_ins_code
_atom_site.Cartn_x
_atom_site.Cartn_y
_atom_site.Cartn_z
_atom_site.occupancy
_atom_site.B_iso_or_equiv
_atom_site.Cartn_x_esd
_atom_site.Cartn_y_esd
_atom_site.Cartn_z_esd
_atom_site.occupancy_esd
_atom_site.B_iso_or_equiv_esd
_atom_site.pdbx_formal_charge
_atom_site.auth_seq_id
_atom_site.auth_comp_id
_atom_site.auth_asym_id
_atom_site.auth_atom_id
_atom_site.pdbx_PDB_model_num
```

ATOM	1	N		TRP	A	1	5	?	8.519	-0.751	10.738	1.00	13.37	?	?	?	?	?	?	5	
ATOM	2	C	CA	.	TRP	A	1	5	?	7.743	-1.668	11.585	1.00	13.42	?	?	?	?	?	?	5
ATOM	3	C	C	.	TRP	A	1	5	?	6.786	-2.502	10.667	1.00	13.47	?	?	?	?	?	?	5
ATOM	4	O	O	.	TRP	A	1	5	?	6.422	-2.085	9.607	1.00	13.57	?	?	?	?	?	?	5
ATOM	5	C	CB	.	TRP	A	1	5	?	6.997	-0.917	12.645	1.00	13.34	?	?	?	?	?	?	5
ATOM	6	C	CG	.	TRP	A	1	5	?	5.784	-0.209	12.221	1.00	13.40	?	?	?	?	?	?	5
ATOM	7	C	CD1	.	TRP	A	1	5	?	5.681	1.084	11.797	1.00	13.29	?	?	?	?	?	?	5
ATOM	8	C	CD2	.	TRP	A	1	5	?	4.417	-0.667	12.221	1.00	13.34	?	?	?	?	?	?	5
ATOM	9	N	NE1	.	TRP	A	1	5	?	4.388	1.418	11.515	1.00	13.30	?	?	?	?	?	?	5
ATOM	10	C	CE2	.	TRP	A	1	5	?	3.588	0.375	11.797	1.00	13.35	?	?	?	?	?	?	5
ATOM	11	C	CE3	.	TRP	A	1	5	?	3.837	-1.877	12.645	1.00	13.39	?	?	?	?	?	?	5
ATOM	12	C	CZ2	.	TRP	A	1	5	?	2.216	0.208	11.656	1.00	13.39	?	?	?	?	?	?	5
ATOM	13	C	CZ3	.	TRP	A	1	5	?	2.465	-2.043	12.504	1.00	13.33	?	?	?	?	?	?	5
ATOM	14	C	CH2	.	TRP	A	1	5	?	1.654	-1.001	12.009	1.00	13.34	?	?	?	?	?	?	5

redundant annotations

inefficient representation

repetitive information

Initial Input: mmCIF or FASTA files

```
;LCB0 - Prolactin precursor - Bovine
; a sample sequence in FASTA format
MDSKGSSQKGSRLLLLLVVSNLLCQGVVSTPVCPNGPGNCQVSLRDLFDRAVMVSHYIHDLSS
EMFNEFDKRYAQGKGFITMALNSCHTSSLPTPEDKEQAQQTHHEVLMMSLILGLLRSWNDPLYHL
VTEVRGMKGAPDAILSRAIEIEEENKRLLEGMEMIFGQVIPGAKETEPYPVWSGLPSLQTKDED
ARYSAFYNLLHCLRRDSSKIDTYLKLLNCRIYNNNC*

>MCHU - Calmodulin - Human, rabbit, bovine, rat, and chicken
MADQLTEEQIAEFKEAFSLFDKDGDTITTKELGTVMRSLGQNPTAEELQDMINEVDADGNGTID
FPEFLTMMARKMKDTDSEEEIREAFRVDKDGNGYISAAELRHVMTNLGEKLTDEEVDEMIREA
DIDGDGQVNYEEFVQMMTAK*

>gi|5524211|gb|AAD44166.1| cytochrome b [Elephas maximus maximus]
LCLYTHIGRNIYYGSYLYSETWNTGIMLLLITMATAFMGYVLPWGQMSFWGATVITNLFSAI PYIGTNLV
EWIWGGFSVDKATLNRRFFAHFILPFTMVALAGVHLTLHETGSNNPLGLTSDSDKIPFHPYYTIKDFLG
LLILLLLLLALLSPDMLGDPDNHMPADPLNTPLHIKPEWYFLFAYAILRSVPNKLGGVLALFLSIVIL
GLMPFLHTSKHRSMMLRPLSQALFWTLTMDLLTWIGSQPVEYPYTIIGQMASILYFSIIILAFLPIAGX
IENY
```

Parsing [1]

- only certain metadata (more from mmCIF)
- change MSE residues into MET

Genetic Search [1]

For MSAs

- JackHMMER
 - MGnify: MSA depth 5,000
 - UniRef90: MSA depth 10,000
- HHBlits
 - Uniclust30 + BFD: MSA depth unlimited
- MSAs duplicated and stacked

flags: JackHMMER: -N 1 -E 0.0001 –incE 0.0001 –F1 0.0005 –F2 0.00005 –F3 0.0000005. HHBlits: -n 3 -e 0.001 -realign_{max} 100000 -maxfilt 100000 -min_{prefilterhits} 1000 -maxseq 100000.

Template Search [1]

- UniRef90 MSA from prior search used for PDB70 search using HHSearch.
- Filter out:
 - released after the input sequence
 - or identical to the input sequence
 - too small
- At inference use top 4 templates

Training Data [1]

- 75:25 self-distillation : known structure (PDB)
- stochastic filters (next)

Filtering [1]

- stochastic filters:
 - Input mmCIFs are restricted to have resolution less than 9 Å. This is not a very restrictive filter and only removes around 0.2% of structures.
 - Longer protein chains are selected with higher probability.
 - Also favour protein chains from smaller clusters. They use 40% sequence identity clusters of the Protein Data Bank clustered with MMSeqs2.
 - Sequences are filtered out when any single amino acid accounts for more than 80% of the input primary sequence. This filter removes about 0.8% of sequences.

MSA block deletion [1]

- MSAs grouped by tool, sorted by block output? (e-value?)
 - similar sequences are likely to be adjacent
 - block deletion tends to remove similarities (ie. whole branch phylogeny)

MSA clustering [1]

- Similarity clusters used to randomly select subset of MSA sequences
 - to reduce computational cost from attention modules, reduce N_{seq}
- K-means, input sequence used as first cluster center
- masking
- hamming distance measure for remaining selections

Residue cropping [1]

During training:

- ① unclamped & clamped - sampling start index from uniform distributions
- ② Cropped with fixed size N_{res}

Featurization and model inputs [1]

- **target_{feat}** This is a feature of size [Nres, 21] consisting of the “aatype” feature.
- **residue_{index}** This is a feature of size [Nres] consisting of the “residue_{index}” feature.
- **msafeat** This is a feature of size [Nclust, Nres, 49] constructed by concatenating “cluster_{msa}”, “cluster_{hasdeletion}”, “cluster_{deletionvalue}”, “cluster_{deletionmean}”, “cluster_{profile}”. We draw $N_{cycle} \times N_{ensemble}$ random samples from this feature to provide each recycling/ensembling iteration of the network with a different sample (see subsubsection 1.11.2).
- **extra_{msafeat}** This is a feature of size [$N_{extra_{seq}}$, Nres, 25] constructed by concatenating “extra_{msa}”, “extra_{msahasdeletion}”, “extra_{msadeletionvalue}”. Together with ‘msafeat’ above we also draw $N_{cycle} \times N_{ensemble}$ random samples from this feature (see subsubsection 1.11.2).

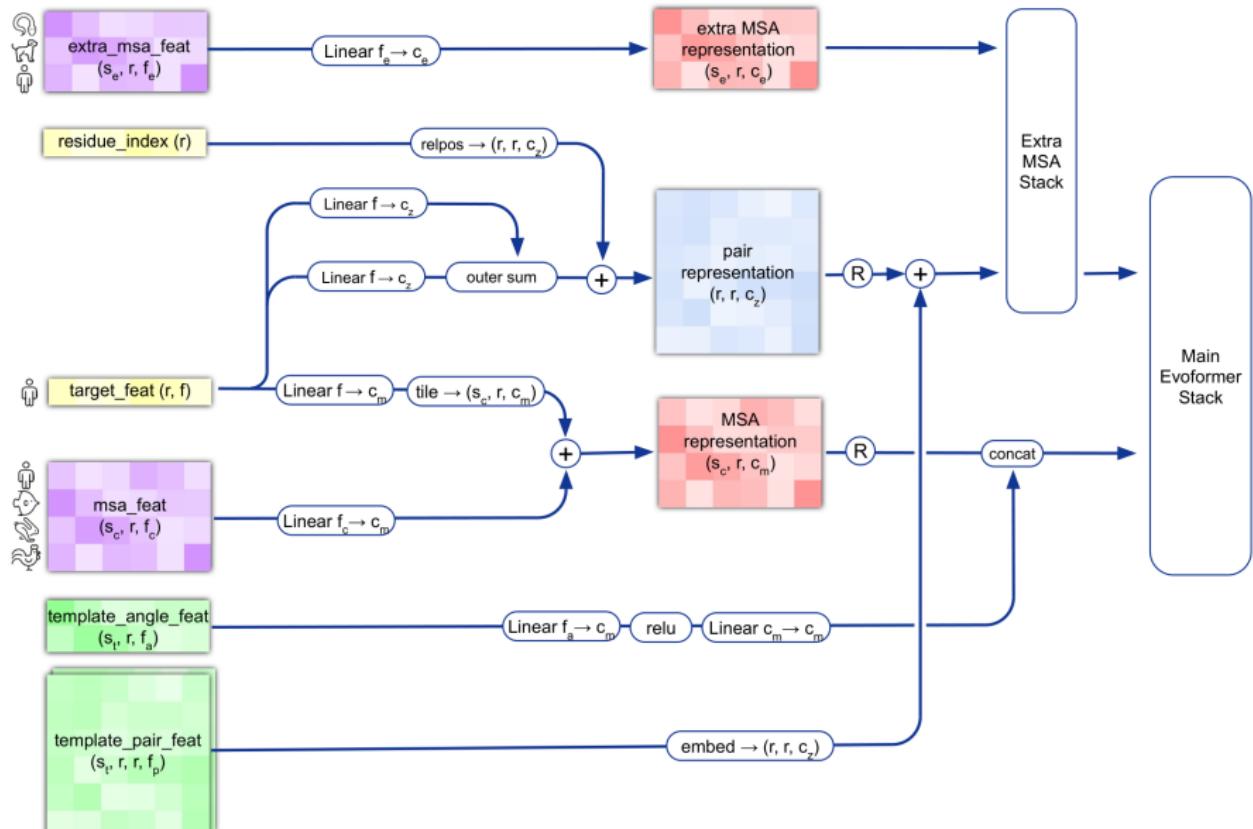
Featurization and model inputs [1]

- **template_{pairfeat}** This is a feature of size [Ntempl, Nres, Nres, 88] and consists of concatenation of the pair residue features “template_{distogram}”, “template_{unitvector}”, and also several residue features, which are transformed into pair features. The “template_{aatype}” feature is included via tiling and stack- ing (this is done twice, in both residue directions). Also the mask features “template_{pseudobetamask}” and “template_{backboneframemask}” are included, where the feature $f_{ij} = \text{mask}_i \cdot \text{mask}_j$. -
template_{anglefeat} This is a feature of size [Ntempl, Nres, 51] constructed by concatenating the following features: “template_{aatype}”, “template_{torsionangles}”, “template_{alttorsionangles}”, and “template_{torsionanglesmask}”.

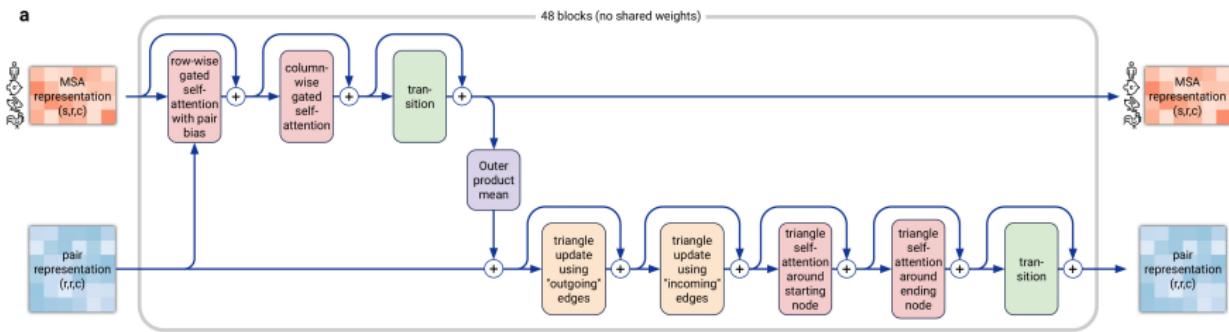
Self-distillation dataset [1]

- Build dataset (on unlabeled sequences):
 - 1 Make MSA for every cluster in Uniclust30
 - 2 Remove sequences that appear in another sequences MSA
 - 3 Keep sequences of $200 < \text{length} < 1024$
 - 4 Remove sequences where MSA < 200 alignments
- For predicted structures:
 - train 'undistilled' model on just PDB dataset
 - use this model to predict above set
 - for every residue pair, computer confidence metric using KL-divergence between distance distribution and a reference distribution
 - reference distribution
- self-distillation training took ~2 weeks

Input embeddings [1]



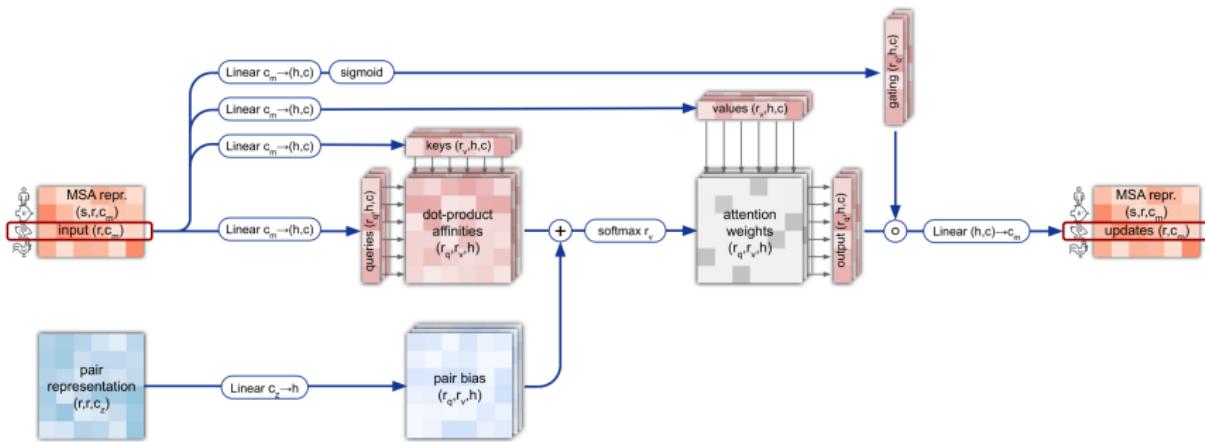
EvoFormer: Overview [1]



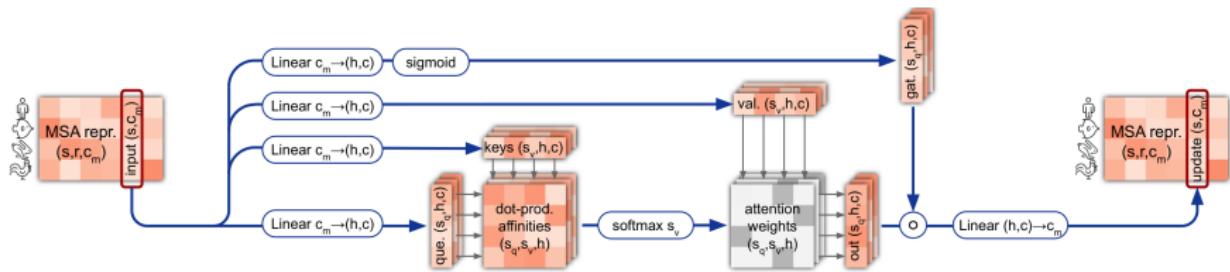
EvoFormer: Overview [1]

- cast as a graph inference problem
- cross-optimization and information flow between MSA representation and pair-wise representation
- layer normalization

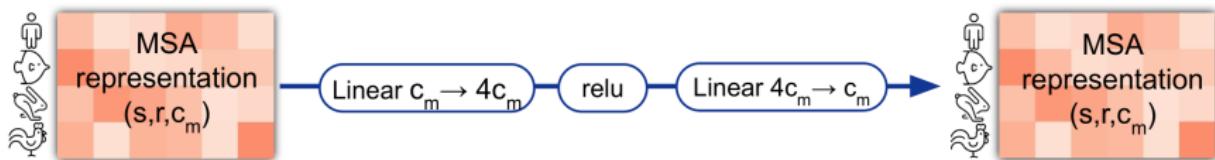
EvoFormer: Row wise Gated Attention [1]



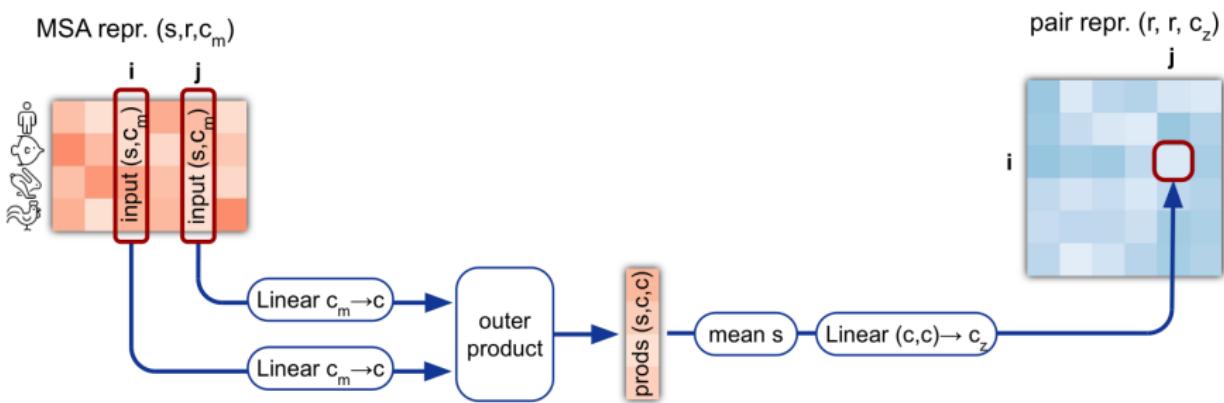
EvoFormer: Column wise Gated Attention [1]



EvoFormer: MSA Translation Layer [1]

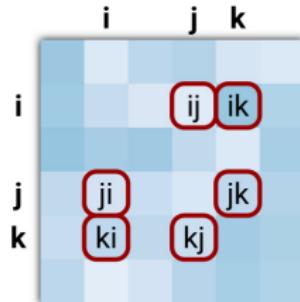


EvoFormer: Outer-Product Mean [1]

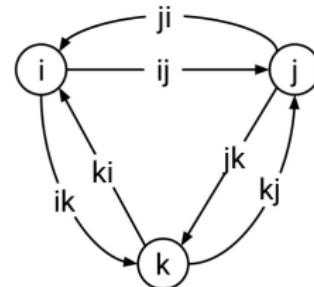


EvoFormer: Residue Pairs [1]

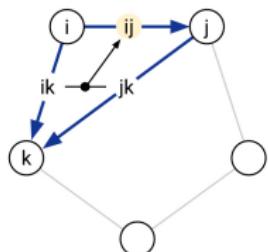
b pair representation
(r, r, c')



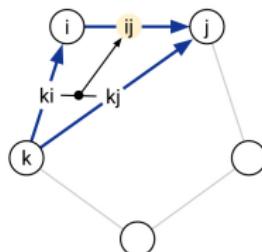
corresponding edges
in a graph



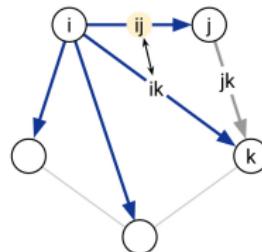
Triangle multiplicative update
using "outgoing" edges



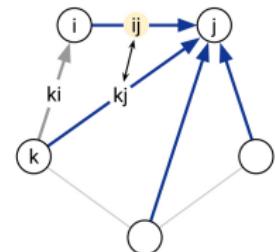
Triangle multiplicative update
using "incoming" edges



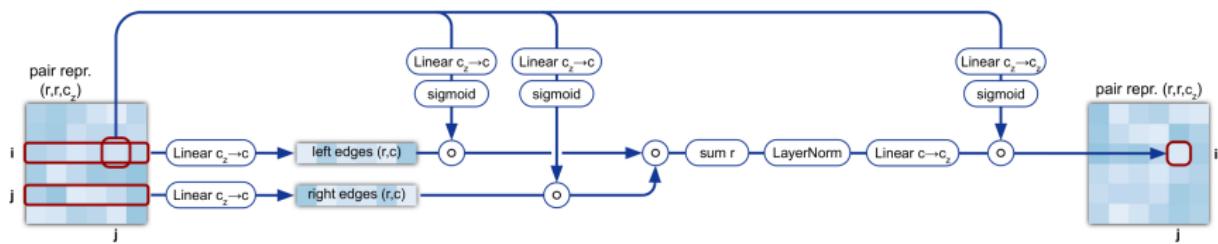
Triangle self-attention around
starting node



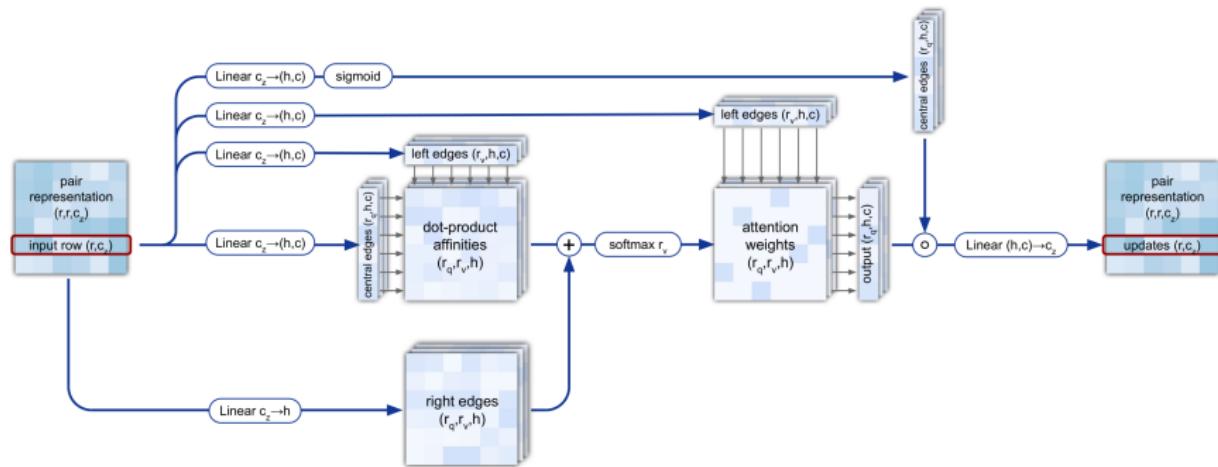
Triangle self-attention around
ending node



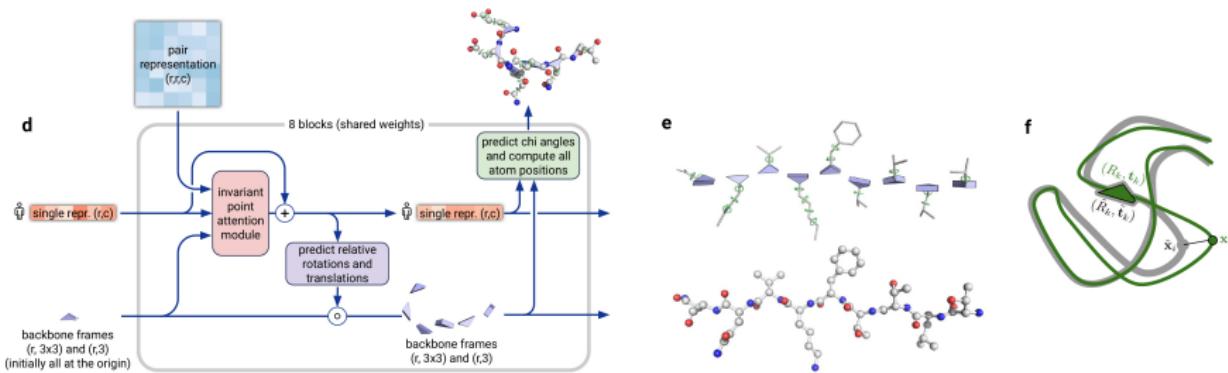
EvoFormer: Triangular Multiplicative Update [1]



EvoFormer: Triangular Self-Attention [1]



Structure Module: Overview [1]



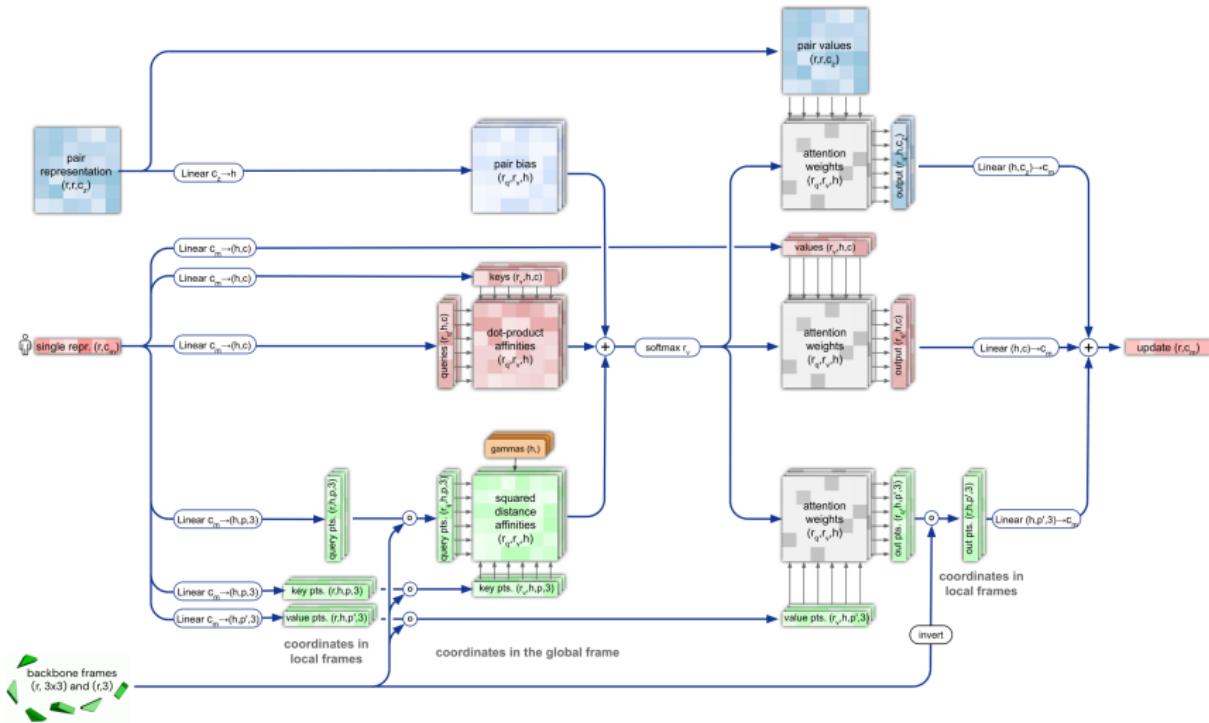
Structure Module: Frame Representation

$$\mathbf{M} = \begin{pmatrix} a_{11} & a_{12} & a_{13} \\ a_{21} & a_{22} & a_{23} \\ a_{31} & a_{32} & a_{33} \\ 0 & 0 & 0 \\ a_{14} \\ a_{24} \\ a_{34} \\ 1 \end{pmatrix}$$

Example transform cite:SpatialTransformationMatrices

- rotation + translation transforms $T_i := (R_i, t_i)$
- no reflection, scaling, or shear

Structure Module: IPA [1]



Structure Module: Algorithm Part 1 [1]

Algorithm 20 Structure module

def StructureModule $\left(\{\mathbf{s}_i^{\text{initial}}\}, \{\mathbf{z}_{ij}\}, N_{\text{layer}} = 8, c = 128, \mathbf{s}_i^{\text{initial}} \in \mathbb{R}^{c_s} \right.$
$$\left. \{T_i^{\text{true},f}\}, \{T_i^{\text{alt truth},f}\}, \{\vec{\alpha}_i^{\text{true},f}\}, \{\vec{\alpha}_i^{\text{alt truth},f}\}, \{\vec{\mathbf{x}}_i^{\text{true},a}\}, \{\vec{\mathbf{x}}_i^{\text{alt truth},a}\} \right) :$$

- 1: $\mathbf{s}_i^{\text{initial}} \leftarrow \text{LayerNorm}(\mathbf{s}_i^{\text{initial}})$
- 2: $\mathbf{z}_{ij} \leftarrow \text{LayerNorm}(\mathbf{z}_{ij})$
- 3: $\mathbf{s}_i = \text{Linear}(\mathbf{s}_i^{\text{initial}}) \quad \mathbf{s}_i \in \mathbb{R}^{c_s}$
- 4: $T_i = (\mathbf{I}, \vec{\mathbf{0}}) \quad \mathbf{I} \in \mathbb{R}^{3 \times 3}, \vec{\mathbf{0}} \in \mathbb{R}^3$

Structure Module: Algorithm Part 2 [1]

```
5: for all  $l \in [1, \dots, N_{\text{layer}}]$  do      # shared weights
6:    $\{\mathbf{s}_i\} += \text{InvariantPointAttention}(\{\mathbf{s}_i\}, \{\mathbf{z}_{ij}\}, \{T_i\})$ 
7:    $\mathbf{s}_i \leftarrow \text{LayerNorm}(\text{Dropout}_{0.1}(\mathbf{s}_i))$ 
# Transition.
8:    $\mathbf{s}_i \leftarrow \mathbf{s}_i + \text{Linear}(\text{relu}(\text{Linear}(\text{relu}(\text{Linear}(\mathbf{s}_i)))))$  all intermediate activations  $\in \mathbb{R}^{c_s}$ 
9:    $\mathbf{s}_i \leftarrow \text{LayerNorm}(\text{Dropout}_{0.1}(\mathbf{s}_i))$ 
# Update backbone.
10:   $T_i \leftarrow T_i \circ \text{BackboneUpdate}(\mathbf{s}_i)$ 
# Predict side chain and backbone torsion angles  $\omega, \phi, \psi, \chi_1, \chi_2, \chi_3, \chi_4$ 
11:   $\mathbf{a}_i = \text{Linear}(\mathbf{s}_i) + \text{Linear}(\mathbf{s}_i^{\text{initial}})$   $\mathbf{a}_i \in \mathbb{R}^c$ 
12:   $\mathbf{a}_i \leftarrow \mathbf{a}_i + \text{Linear}(\text{relu}(\text{Linear}(\text{relu}(\mathbf{a}_i))))$  all intermediate activations  $\in \mathbb{R}^c$ 
13:   $\mathbf{a}_i \leftarrow \mathbf{a}_i + \text{Linear}(\text{relu}(\text{Linear}(\text{relu}(\mathbf{a}_i))))$  all intermediate activations  $\in \mathbb{R}^c$ 
14:   $\vec{\alpha}_i^f = \text{Linear}(\text{relu}(\mathbf{a}_i))$   $\vec{\alpha}_i^f \in \mathbb{R}^2, f \in \mathcal{S}_{\text{torsion names}}$ 
# Auxiliary losses in every iteration.
15:   $(R_i, \vec{\mathbf{t}}_i) = T_i$ 
16:   $\vec{\mathbf{x}}_i^{\text{C}\alpha} = \vec{\mathbf{t}}_i$ 
17:   $\mathcal{L}_{\text{aux}}^l = \left( \text{computeFAPE}(\{T_i\}, \{\vec{\mathbf{x}}_i^{\text{C}\alpha}\}, \{T_i^{\text{true}}\}, \{\vec{\mathbf{x}}_i^{\text{true,C}\alpha}\}, \epsilon = 10^{-12}) \right.$ 
18:     $\left. + \text{torsionAngleLoss}(\{\vec{\alpha}_i^f\}, \{\vec{\alpha}_i^{\text{true},f}\}, \{\vec{\alpha}_i^{\text{alt truth},f}\}) \right)$ 
```

Structure Module: Algorithm Part 3 [1]

```
23:  $\mathcal{L}_{\text{aux}} = \text{mean}_l(\{\mathcal{L}_{\text{aux}}^l\})$ 
24:  $T_i^f, \vec{x}_i^a = \text{computeAllAtomCoordinates}(T_i, \vec{\alpha}_i^f)$   $a \in \mathcal{S}_{\text{atom names}}$ 
25:  $T_i^f \leftarrow \text{concat}(T_i, T_i^f)$ 

# Final loss on all atom coordinates.
26:  $\{T_i^{\text{true}, f}\}, \{\vec{x}_i^{\text{true}, a}\} \leftarrow \text{renameSymmetricGroundTruthAtoms}($ 
27:  $\{T_i^f\}, \{\vec{x}_i^a\}, \{T_i^{\text{true}, f}\}, \{T_i^{\text{alt truth}, f}\}, \{\vec{x}_i^{\text{true}, a}\}, \{\vec{x}_i^{\text{alt truth}, a}\})$ 
28:  $\mathcal{L}_{\text{FAPE}} = \text{computeFAPE}(\{T_i^f\}, \{\vec{x}_i^a\}, \{T_i^{\text{true}, f}\}, \{\vec{x}_i^{\text{true}, a}\}, \epsilon = 10^{-4})$ 

# Predict model confidence.
29:  $\{r_i^{\text{true LDDT}}\} = \text{perResidueLDDT_Ca}(\{\vec{x}_i^a\}, \{\vec{x}_i^{\text{true}, a}\})$ 
30:  $\{r_i^{\text{pLDDT}}\}, \mathcal{L}_{\text{conf}} = \text{predictPerResidueLDDT_Ca}(\{\mathbf{s}_i\}, \{r_i^{\text{true LDDT}}\})$ 
31: return  $\{\vec{x}_i^a\}, \{r_i^{\text{pLDDT}}\}, \mathcal{L}_{\text{FAPE}}, \mathcal{L}_{\text{conf}}, \mathcal{L}_{\text{aux}}$ 
```

Loss Functions [1]

$$\mathcal{L} = \begin{cases} 0.5\mathcal{L}_{\text{FAPE}} + 0.5\mathcal{L}_{\text{aux}} + 0.3\mathcal{L}_{\text{dist}} + 2.0\mathcal{L}_{\text{msa}} + 0.01\mathcal{L}_{\text{conf}} & \text{training} \\ 0.5\mathcal{L}_{\text{FAPE}} + 0.5\mathcal{L}_{\text{aux}} + 0.3\mathcal{L}_{\text{dist}} + 2.0\mathcal{L}_{\text{msa}} + 0.01\mathcal{L}_{\text{conf}} + 0.01\mathcal{L}_{\text{exp resolved}} + 1.0\mathcal{L}_{\text{viol}} & \text{fine-tuning} \end{cases}, \quad (7)$$

- weighted sum
- weighted to reduce importance of short sequences

Loss Functions & Auxillary Heads [1]

- 1 Side chain and backbone torsion angle loss
- 2 Frame aligned point error (FAPE)
 - Configurations with $\text{FAPE}(X,Y) = 0$
 - Metric properties of FAPE
- 3 Chiral properties of AlphaFold and its loss
- 4 Model confidence prediction (pLDDT)
- 5 TM-score prediction
- 6 Distogram prediction
- 7 Masked MSA prediction
- 8 "Experimentally resolved" prediction
- 9 Structural violations

Loss Functions: FAPE

Algorithm 28 Compute the Frame aligned point error

def computeFAPE($\{T_i\}, \{\vec{x}_j\}, \{T_i^{\text{true}}\}, \{\vec{x}_j^{\text{true}}\}, Z = 10\text{\AA}, d_{\text{clamp}} = 10\text{\AA}, \epsilon = 10^{-4}\text{\AA}^2$) :

$$T_i, T_i^{\text{true}} \in (\mathbb{R}^{3 \times 3}, \mathbb{R}^3)$$
$$\vec{x}_j, \vec{x}_j^{\text{true}} \in \mathbb{R}^3,$$
$$i \in \{1, \dots, N_{\text{frames}}\}, j \in \{1, \dots, N_{\text{atoms}}\}$$

1: $\vec{x}_{ij} = T_i^{-1} \circ \vec{x}_j \quad \vec{x}_{ij} \in \mathbb{R}^3$

2: $\vec{x}_{ij}^{\text{true}} = T_i^{\text{true}-1} \circ \vec{x}_j^{\text{true}} \quad \vec{x}_{ij}^{\text{true}} \in \mathbb{R}^3$

3: $d_{ij} = \sqrt{\|\vec{x}_{ij} - \vec{x}_{ij}^{\text{true}}\|^2 + \epsilon} \quad d_{ij} \in \mathbb{R}$

4: $\mathcal{L}_{\text{FAPE}} = \frac{1}{Z} \text{mean}_{i,j}(\min(d_{\text{clamp}}, d_{ij}))$

5: **return** $\mathcal{L}_{\text{FAPE}}$

[1]

- Variation of commonly used root-mean-squared deviation (RMSD) of atomic positions
- not invariant to reflections, preventing proteins of the wrong chirality. [3], [1]

AlphaFold Inference [1]

- AlphaFold receives input features derived from:
 - the amino-acid sequence
 - MSA
 - templates (see subsubsection 1.2.9)
- outputs features:
 - atom coordinates
 - the distogram
 - per-residue confidence scores.
- Recycling x3
 - initial recycled inputs are zero

Algorithm 2 outlines the main steps (see also Fig 1e and the corresponding description in the main article).

AlphaFold Training [1]

Table 4 | AlphaFold training protocol. We train each stage until convergence with the approximate timings and number of samples provided.

Model	Initial training	Fine-tuning
Number of templates N_{templ}	4	4
Sequence crop size N_{res}	256	384
Number of sequences N_{seq}	128	512
Number of extra sequences $N_{\text{extra_seq}}$	1024	5120
Parameters initialized from	Random	Initial training
Initial learning rate	10^{-3}	$5 \cdot 10^{-4}$
Learning rate linear warm-up samples	128000	0
Structural violation loss weight	0.0	1.0
Training samples ($\cdot 10^6$)	≈ 10	≈ 1.5
Training time	≈ 7 days	≈ 4 days

Results [1]

They did well

Results [1]

They did well

Results: Positional Encodings [1]

Novel Folds

They did well

Ablation Studies [1]

Baseline for all ablation models: Full model without noisy-student self-attention
Ablations:

- 1 With noisy-student self-distillation training
- 2 No templates
- 3 No raw MSA (use MSA pairwise frequencies)
- 4 No triangles, biasing, or gating (use axial attention)
- 5 No recycling
- 6 No IPA (use direct projection)
- 7 No invariant IPA & no recycling
- 8 No end-to-end structure gradients (keep auxiliary heads)
- 9 No auxiliary distogram head
- 10 No auxiliary masked MSA head

Network Probing [1]

todo

Attention Visualization [1]

todo

Algorithm 24 Compute all atom coordinates [1]

Algorithm 24 Compute all atom coordinates

def computeAllAtomCoordinates($T_i, \vec{\alpha}_i^f, F_i^{\text{aatype}}$) :

1: $\hat{\vec{\alpha}}_i^f = \vec{\alpha}_i^f / \|\vec{\alpha}_i^f\|$

2: $(\vec{\omega}_i, \vec{\phi}_i, \vec{\psi}_i, \vec{\chi}_{1i}, \vec{\chi}_{2i}, \vec{\chi}_{3i}, \vec{\chi}_{4i}) = \hat{\vec{\alpha}}_i^f$

Make extra backbone frames.

3: $r_i = F_i^{\text{aatype}}$

4: $T_{i1} = T_i \circ T_{r_i, (\omega \rightarrow \text{bb})}^{\text{lit}} \circ \text{makeRotX}(\vec{\omega}_i)$

5: $T_{i2} = T_i \circ T_{r_i, (\phi \rightarrow \text{bb})}^{\text{lit}} \circ \text{makeRotX}(\vec{\phi}_i)$

6: $T_{i3} = T_i \circ T_{r_i, (\psi \rightarrow \text{bb})}^{\text{lit}} \circ \text{makeRotX}(\vec{\psi}_i)$

Make side chain frames (chain them up along the side chain).

7: $T_{i4} = T_i \circ T_{r_i, (\chi_1 \rightarrow \text{bb})}^{\text{lit}} \circ \text{makeRotX}(\vec{\chi}_{1i})$

8: $T_{i5} = T_{i4} \circ T_{r_i, (\chi_2 \rightarrow \chi_1)}^{\text{lit}} \circ \text{makeRotX}(\vec{\chi}_{2i})$

9: $T_{i6} = T_{i5} \circ T_{r_i, (\chi_3 \rightarrow \chi_2)}^{\text{lit}} \circ \text{makeRotX}(\vec{\chi}_{3i})$

10: $T_{i7} = T_{i6} \circ T_{r_i, (\chi_4 \rightarrow \chi_3)}^{\text{lit}} \circ \text{makeRotX}(\vec{\chi}_{4i})$

Map atom literature positions to the global frame.

11: $\vec{\mathbf{x}}^a = \text{concat}_{\vec{\mathbf{x}}^f, \vec{\mathbf{x}}^{\text{lit}}} \left(\{T_i^f \circ \vec{\mathbf{x}}^{\text{lit}}_{\cdot, \cdot, \cdot, \cdot}\} \right)$

Algorithm 23 Backbone update [1]

Algorithm 23 Backbone update

def BackboneUpdate(\mathbf{s}_i) :

1: $b_i, c_i, d_i, \vec{\mathbf{t}}_i = \text{Linear}(\mathbf{s}_i)$

$b_i, c_i, d_i \in \mathbb{R}, \vec{\mathbf{t}}_i \in \mathbb{R}^3$

Convert (non-unit) quaternion to rotation matrix.

2: $(a_i, b_i, c_i, d_i) \leftarrow (1, b_i, c_i, d_i) / \sqrt{1 + b_i^2 + c_i^2 + d_i^2}$

3: $R_i = \begin{pmatrix} a_i^2 + b_i^2 - c_i^2 - d_i^2 & 2b_i c_i - 2a_i d_i & 2b_i d_i + 2a_i c_i \\ 2b_i c_i + 2a_i d_i & a_i^2 - b_i^2 + c_i^2 - d_i^2 & 2c_i d_i - 2a_i b_i \\ 2b_i d_i - 2a_i c_i & 2c_i d_i + 2a_i b_i & a_i^2 - b_i^2 - c_i^2 + d_i^2 \end{pmatrix}$

4: $T_i = (R_i, \vec{\mathbf{t}}_i)$

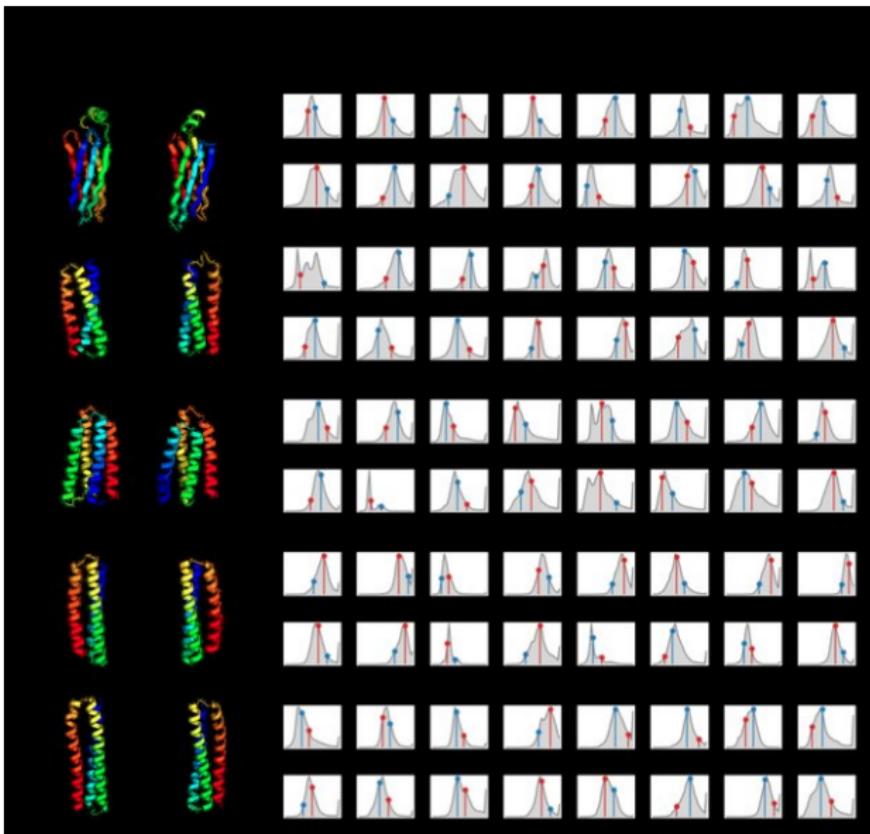
5: **return** T_i

Algorithm 29 Predict model confidence pLDDT [1]

Algorithm 29 Predict model confidence pLDDT

```
def predictPerResidueLDDT_Ca({si}, vbins = [1, 3, 5, ..., 99]T, {ritrue LDDT}, c = 128) :
    1: ai = relu(Linear(relu(Linear(LayerNorm(si)))))
        ai, and intermediate activations ∈ Rc
    2: pipLDDT = softmax(Linear(ai))
        pipLDDT ∈ R|vbins|
    3: pitrue LDDT = one_hot(ritrue LDDT, vbins)
    4: Lconf = meani(pitrue LDDTT log pipLDDT)
    5: ripLDDT = pipLDDTT vbins
        ripLDDT ∈ R
    6: return {ripLDDT}, Lconf
```

Distograms [1]



Algorithm 31 Generic recycling training procedure [1]

Algorithm 31 Generic recycling training procedure

```
def RecyclingTraining( $\{\text{inputs}_c\}$ ,  $N_{\text{cycle}}$ ) :  
    1:  $N' = \text{uniform}(1, N_{\text{cycle}})$       # shared value across the batch  
    2: outputs = 0  
    # Recycling iterations  
    3: for all  $c \in [1, \dots, N']$  do      # shared weights  
        4: outputs  $\leftarrow \text{stopgrad}(\text{outputs})$       # no gradients between iterations  
        5: outputs  $\leftarrow \text{Model}(\text{inputs}_c, \text{outputs})$   
    6: end for  
    7: return loss(outputs)      # only the final iteration's outputs are used
```

Algorithm 30 Generic recycling inference procedure [1]

Algorithm 30 Generic recycling inference procedure

```
def RecyclingInference( {inputsc} , Ncycle ) :  
    1: outputs = 0  
    # Recycling iterations  
    2: for all c ∈ [1, . . . , Ncycle] do      # shared weights  
    3:     outputs ← Model(inputsc, outputs)  
    4: end for  
    5: return outputs
```

Algorithm 32 Embedding of evoformer and structure module outputs for recycling [1]

Algorithm 32 Embedding of Evoformer and Structure module outputs for recycling

def RecyclingEmbedder($\{\mathbf{m}_{1i}\}$, $\{\mathbf{z}_{ij}\}$, $\{\vec{\mathbf{x}}_i^{\text{C}^\beta}\}$) :

Embed pair distances of backbone atoms:

1: $d_{ij} = \left\| \vec{\mathbf{x}}_i^{\text{C}^\beta} - \vec{\mathbf{x}}_j^{\text{C}^\beta} \right\|$ C^α used for glycine

2: $\mathbf{d}_{ij} = \text{Linear}(\text{one_hot}(d_{ij}, \mathbf{v}_{\text{bins}} = [3\frac{1}{8}\text{\AA}, 5\frac{1}{8}\text{\AA}, \dots, 21\frac{1}{8}\text{\AA}]))$ $\mathbf{d}_{ij} \in \mathbb{R}^{cz}$

Embed output Evoformer representations:

3: $\tilde{\mathbf{z}}_{ij} = \mathbf{d}_{ij} + \text{LayerNorm}(\mathbf{z}_{ij})$

4: $\tilde{\mathbf{m}}_{1i} = \text{LayerNorm}(\mathbf{m}_{1i})$

5: **return** $\{\tilde{\mathbf{m}}_{1i}\}, \{\tilde{\mathbf{z}}_{ij}\}$

Algorithm 26 Rename symmetric ground truth atoms [1]

Algorithm 26 Rename symmetric ground truth atoms

def renameSymmetricGroundTruthAtoms($\{T_i^f\}$, $\{\vec{x}_i^a\}$, $\{T_i^{\text{true},f}\}$, $\{T_i^{\text{alt truth},f}\}$, $\{\vec{x}_i^{\text{true},a}\}$, $\{\vec{x}_i^{\text{alt truth},a}\}$) :

- 1: $d_{(i,a),(j,b)} = \left\| \vec{x}_i^a - \vec{x}_j^b \right\| \quad \forall (j, b) \in \mathcal{S}_{\text{non-ambiguous atoms}}$
- 2: $d_{(i,a),(j,b)}^{\text{true}} = \left\| \vec{x}_i^{\text{true},a} - \vec{x}_j^{\text{true},b} \right\| \quad \forall (j, b) \in \mathcal{S}_{\text{non-ambiguous atoms}}$
- 3: $d_{(i,a),(j,b)}^{\text{alt truth}} = \left\| \vec{x}_i^{\text{alt truth},a} - \vec{x}_j^{\text{true},b} \right\| \quad \forall (j, b) \in \mathcal{S}_{\text{non-ambiguous atoms}}$
- 4: **for all** $i \in [1 \dots N_{\text{res}}]$ **do**
- 5: **if** $\text{sum}_{a,(j,b)} \left(\left| d_{(i,a),(j,b)} - d_{(i,a),(j,b)}^{\text{alt truth}} \right| \right) < \text{sum}_{a,(j,b)} \left(\left| d_{(i,a),(j,b)} - d_{(i,a),(j,b)}^{\text{true}} \right| \right)$ **then**
- 6: $\vec{x}_i^{\text{true},a} \leftarrow \vec{x}_i^{\text{alt truth},a}$
- 7: $T_i^{\text{true},f} \leftarrow T_i^{\text{alt truth},f}$
- 8: **end if**
- 9: **end for**
- 10: **return** $\{T_i^{\text{true},f}\}$, $\{\vec{x}_i^{\text{true},a}\}$

Algorithm 27 Side chain and backbone torsion angle loss [1]

Algorithm 27 Side chain and backbone torsion angle loss

```
def torsionAngleLoss({ $\vec{\alpha}_i^f$ }, { $\vec{\alpha}_i^{\text{true}, f}$ }, { $\vec{\alpha}_i^{\text{alt truth}, f}$ }) :  
    1:  $\ell_i^f = \|\vec{\alpha}_i^f\|$   
    2:  $\hat{\vec{\alpha}}_i^f = \vec{\alpha}_i^f / \ell_i^f$   
    3:  $\mathcal{L}_{\text{torsion}} = \text{mean}_{i,f}(\text{minimum}(\|\hat{\vec{\alpha}}_i^f - \vec{\alpha}_i^{\text{true}, f}\|^2, \|\hat{\vec{\alpha}}_i^f - \vec{\alpha}_i^{\text{alt truth}, f}\|^2))$   
    4:  $\mathcal{L}_{\text{anglenorm}} = \text{mean}_{i,f}(|\ell_i^f - 1|)$   
    5: return  $\mathcal{L}_{\text{torsion}} + 0.02\mathcal{L}_{\text{anglenorm}}$ 
```

Algorithm 25 Make a transformation that rotates around the x-axis [1]

Algorithm 25 Make a transformation that rotates around the x-axis

def makeRotX($\vec{\alpha}$) : $\vec{\alpha} \in R^2$ with $\|\vec{\alpha}\| = 1$

1: $R = \begin{pmatrix} 1 & 0 & 0 \\ 0 & \alpha_1 & -\alpha_2 \\ 0 & \alpha_2 & \alpha_1 \end{pmatrix}$

2: $\vec{t} = \begin{pmatrix} 0 \\ 0 \\ 0 \end{pmatrix}$

3: $T = (R, \vec{t})$

4: **return** T

Bibliography I

- [1] John Jumper et al. "Highly Accurate Protein Structure Prediction with AlphaFold". In: *Nature* (July 15, 2021), pp. 1–11. ISSN: 1476-4687. DOI: 10/gk7nfp. URL: <https://www.nature.com/articles/s41586-021-03819-2> (visited on 07/15/2021).
- [2] *PDB101: Learn: Guide to Understanding PDB Data: Beginner's Guide to PDB Structures and the PDBx/mmCIF Format.* RCSB: PDB-101. URL: <https://pdb101.rcsb.org/learn/guide-to-understanding-pdb-data/beginner%80%99s-guide-to-pdb-structures-and-the-pdbx-mm cif-format> (visited on 07/24/2021).
- [3] Carlos Outeiral Rubiera. *AlphaFold 2 Is Here: What's behind the Structure Prediction Miracle* | Oxford Protein Informatics Group. URL: <https://www.blopig.com/blog/2021/07/alphafold-2-is-here-whats-behind-the-structure-prediction-miracle/> (visited on 07/24/2021).